

AI JOB APPLICATION PLATFORM

Product, Workflow & Multi-AI Development Specification

Version 1.2

For Product, Engineering, AI/ML, UX/UI, Backend, DevOps and QA Teams

Core product direction: Build an AI Job Application Manager — not merely a job board or resume writer. The platform should combine verified candidate intelligence, job intelligence, multi-AI reasoning, application assistance and outcome learning.

Executive Overview

The platform is an AI-powered job discovery, resume optimization and application management system designed to help users move from career profile creation through job discovery, application preparation, submission, tracking and interview readiness.

The application must use a multi-AI architecture from the beginning. Different AI models should be assigned to different functions, while deterministic code handles tasks that do not require generative reasoning. Critical outputs should be independently reviewed and validated before reaching the user.

Career Profile → Job Discovery → Job Matching → Job Analysis → Resume Optimization →
Application Preparation → Submission → Tracking → Interview Preparation → Outcome Learning

Development Principles at a Glance

- Use a provider-independent AI Orchestrator rather than hard-coding one AI vendor.
- Store candidate facts in a structured Fact Ledger and keep generated resumes separate from source-of-truth data.
- Use multi-AI produce/review patterns for high-impact outputs.
- Use deterministic rules for dates, salary, eligibility, thresholds, status and scoring logic.
- Require structured JSON between AI pipeline stages wherever possible.
- Do not allow AI to fabricate candidate qualifications, achievements, employment history or credentials.
- Keep the user in control of sensitive declarations and application answers.
- Version prompts, models and AI outputs for debugging, rollback and quality assurance.

1. Product Vision

The platform should act as an AI Job Application Manager. It should understand the user's professional history once and continuously use that verified information across job discovery, matching, resume tailoring, application preparation, tracking and interview preparation.

Positioning: Find it. Match it. Tailor it. Apply. Track it. Get hired.

2. Core Product Principle

Optimize for application quality rather than application volume.

- Determine whether the candidate genuinely fits the role.
- Identify which experience and skills are most relevant.
- Identify missing or weak qualifications.
- Improve presentation without inventing facts.
- Prepare application materials for a specific role.
- Require review whenever the system is uncertain.

3. Core User Journey

The preferred onboarding should minimize registration friction by using the candidate's resume as the primary input for creating the account and Master Career Profile.

Resume Upload → AI Parsing → Draft Account/Profile → Identity & Consent Confirmation → Missing Preferences → Account Activation → Job Matching

The resume should function as the user's initial registration form. The platform should extract available personal and professional data, then ask only for information that is missing, outdated, sensitive, or requires explicit user confirmation.

4. Resume-First AI Account Creation and Onboarding

The default signup path should begin with resume upload rather than a traditional registration form.

Primary onboarding flow

Landing Page → Upload Resume → AI Resume Parsing → Draft Candidate Profile → Confirm Name/Email/Phone → Accept Terms & Privacy → Verify Email/Phone → Ask Only Missing Preferences → Activate Account → Start Job Matching

The system may extract the following from the resume when present:

- Name
- Email address
- Phone number
- City / state / country
- Employment history
- Job titles and employers
- Education and certifications
- Skills, tools and technologies
- Years and level of experience

- Possible target-role categories

Important account rule

Resume-derived identity data should create a draft account/profile, but the account should not become fully active until the user confirms ownership and provides explicit consent.

Reasons include outdated contact information, resumes uploaded by another person, missing identity fields, and the need to capture consent to Terms, Privacy Policy and AI processing.

Identity confirmation

If an email address is found on the resume, display it and ask the user to confirm that it is still valid. Use email OTP or a magic link rather than requiring a password during initial onboarding.

If no valid email exists, ask the user to provide one. Phone verification may be added where required by product or fraud controls.

Dynamic missing-data interview

Do not ask the user to re-enter information already extracted with high confidence. After parsing the resume, the AI should generate only the missing questions required to complete job matching and application preferences.

Examples of questions that may still be required:

- What minimum salary are you targeting?
- Are you open to remote, hybrid or on-site work?
- Which locations are acceptable?
- What employment type do you prefer?
- Are you authorized to work in the target country, and do you require sponsorship?
- Which job titles are you actively targeting?

Activation objective

The user should experience meaningful value immediately after resume upload. Once the draft profile is created, the platform should begin matching jobs and show an initial career/profile summary while the user completes only the necessary confirmations.

Example activation screen:

```
Your career profile is ready.  
Senior Network Engineer | 8+ years experience | Houston, TX  
Cisco • Fortinet • AWS • Network Security  
We found relevant opportunities based on your background.  
Complete the remaining required details to view your strongest matches.
```

5. Acquisition and Activation Flow

The acquisition experience should reinforce the resume-first onboarding model.

```
Traffic / Referral / Search / Social → Landing Page → Upload Resume → AI Builds Draft Profile  
→ User Confirms Identity & Consent → AI Finds Matching Jobs → First Tailored Application →  
Paid Conversion
```

The key activation event is not account creation. It is the point at which the user sees a meaningful job match and has a tailored application or resume ready for review.

Primary conversion hooks

- Resume Upscale — improve a weak resume using verified facts.
- Job Match — show how well the user's actual experience fits a specific job.
- Application Ready — prepare a tailored resume and application package.

6. Master Career Profile

Every user should have a structured Master Career Profile. It becomes the authoritative candidate profile used across the platform and must be separate from generated resume versions.

The profile may be created through:

- Existing resume upload.
- LinkedIn or other professional profile import where permitted.
- Guided interview.
- Manual entry.
- Certificates and supporting documents.

Core profile structure:

```
Candidate
Personal Details
Professional Summary
Employment History
Job Titles and Employers
Employment Dates
Responsibilities
Verified Achievements
Projects
Education
Certifications
Technical Skills
Soft Skills
Management Experience
Industry Experience
Tools / Platforms
Languages
Work Authorization
Geographic Preferences
Salary Preferences
Remote Preferences
Career Goals
```

7. Candidate Fact Ledger

Behind the Master Career Profile, maintain a structured Candidate Fact Ledger. Every high-impact AI-generated claim should be traceable to one or more verified facts.

```
{
  "fact_id": "EXP-001",
  "category": "employment",
  "employer": "ABC Technologies",
  "job_title": "Network Engineer",
  "start_date": "2021-04",
  "end_date": "2024-03",
  "skills": ["Cisco", "Fortinet", "AWS"],
  "responsibilities": ["Managed enterprise networking infrastructure"],
  "verified_by_user": true,
```

```
"source": "uploaded_resume"
}
```

8. Non-Fabrication Rule

AI may rewrite, reorganize, improve, condense and emphasize verified candidate information.

AI must not:

- Create employers or job history.
- Invent certifications or education.
- Change employment dates without confirmation.
- Add technologies or skills the user has not demonstrated.
- Invent responsibilities, performance figures or achievements.

When stronger information is needed, the system should ask the candidate a targeted clarification question and only add the answer to the Fact Ledger after confirmation.

9. Multi-AI Architecture

```
USER
↓
APPLICATION
↓
AI ORCHESTRATOR
↓
TASK CLASSIFICATION
↓
SPECIALIZED AI PIPELINE
↓
VALIDATION / RULE ENGINE
↓
FACT / POLICY CHECK
↓
PRESENTATION OR ACTION
```

No single AI model should be responsible for the entire workflow. The orchestrator should route each task to the model or deterministic service best suited for that function.

10. AI Provider Abstraction

```
AIProvider
├── OpenAIProvider
├── AnthropicProvider
├── GoogleProvider
└── FutureProvider
```

The abstraction layer should allow the platform to:

- Assign providers by task.
- Change models without rebuilding the product.
- Use inexpensive models for extraction and classification.
- Use stronger reasoning models only when required.
- Run independent reviewers from a second provider.
- Fail over when a provider is unavailable.
- Track cost and latency by model/function.

11. AI Model Routing and Configuration

Model routing should be configurable and versioned rather than embedded throughout application code.

```
{
  "career_profile_extraction": {
    "primary": "Provider-A-Fast",
    "review": "Provider-B-Reasoning"
  },
  "resume_generation": {
    "primary": "Provider-A-Writing",
    "review": "Provider-B-Reasoning"
  },
  "job_fit_analysis": {
    "primary": "Provider-B-Reasoning",
    "secondary": "Provider-A-Reasoning"
  }
}
```

12. AI Pipeline 1 — Career Profile Extraction

```
Resume Upload
  ↓
Document Parser
  ↓
AI Extractor
  ↓
Structured Career JSON
  ↓
AI Fact Reviewer
  ↓
User Confirmation
  ↓
Master Career Profile
```

AI 1 — Resume Extractor

- Extract employers, dates and job titles.
- Extract skills, education and certifications.
- Extract responsibilities and stated achievements.
- Return strict structured JSON.

AI 2 — Career Fact Reviewer

- Detect contradictory employment dates.
- Identify missing or ambiguous information.
- Detect duplicated experience.
- Flag suspicious inference.
- Generate clarification questions.

13. AI Pipeline 2 — Resume Upscale

Feature name: Resume Upscale

```
Original Resume
  ↓
Resume Diagnostic AI
  ↓
```

```
Gap Identification
↓
Candidate Questions
↓
Verified Answers
↓
Fact Ledger Update
↓
Resume Writer AI
↓
Resume Verification AI
↓
Final Resume
```

Resume Diagnostic AI

- Detect vague statements and weak descriptions.
- Identify missing metrics, scope or technologies.
- Detect excessive repetition and poor structure.
- Identify ATS issues and inconsistent dates.
- Identify career gaps or unclear progression.

Resume Writer AI

- Write professional summaries.
- Rewrite experience bullets.
- Organize skills and achievements.
- Adjust emphasis for target roles.
- Preserve ATS readability.
- Use only verified profile facts.

Resume Verification AI

A separate reviewer validates every generated factual claim against the Candidate Fact Ledger.

```
{
  "status": "FAIL",
  "unsupported_claims": [
    {
      "statement": "Managed a $10M technology portfolio",
      "reason": "No verified portfolio value exists."
    }
  ]
}
```

14. AI Pipeline 3 — Job Ingestion and Normalization

Jobs may enter the platform from:

- Employer career sites.
- Public ATS feeds.
- Licensed job APIs and providers.
- Employer submissions.
- Browser extension capture.
- User-provided job URLs.

Normalize every job into a common schema:

```
{
  "job_id": "",
  "company": "",
  "title": "",
  "location": "",
  "remote_type": "",
  "employment_type": "",
  "salary_min": null,
  "salary_max": null,
  "job_description": "",
  "requirements": [],
  "preferred_requirements": [],
  "skills": [],
  "education": [],
  "experience_years": null,
  "work_authorization": "",
  "source_url": "",
  "source_type": "",
  "date_detected": "",
  "date_verified": ""
}
```

15. AI Pipeline 4 — Job Requirement Extraction

A dedicated AI should extract and normalize requirements before matching. The matching engine should not depend only on raw job-description prose.

```
{
  "mandatory": [
    "5+ years network engineering",
    "Cisco routing",
    "Firewall administration"
  ],
  "preferred": [
    "AWS",
    "Azure",
    "Kubernetes"
  ],
  "management": false,
  "minimum_education": "Bachelor's preferred"
}
```

16. AI Pipeline 5 — Job Matching

Job matching should combine three layers:

A. Deterministic Matching

- Job-title similarity.
- Location and work arrangement.
- Salary compatibility.
- Work authorization.
- Experience level.
- Required certifications.
- Required skills.

B. Semantic Matching

Use embeddings or equivalent semantic similarity to detect relevant experience expressed in different language.

C. AI Reasoning

Use a reasoning model to interpret contextual fit, transferable experience and meaningful gaps.

Example Match Score

Factor	Weight	Example
Core Experience	25%	24/25
Required Skills	25%	21/25
Role Similarity	15%	14/15
Industry Experience	10%	8/10
Education / Certification	10%	9/10
Location	5%	5/5
Salary	5%	4/5
Preferred Skills	5%	2/5


Example total: 87% match.

17. AI Match Analyst

The AI Match Analyst explains the structured score and identifies strengths, gaps and risk factors. Status labels should be based on configured score thresholds rather than free-form AI opinion.

STRONG MATCH
GOOD MATCH
POSSIBLE MATCH
WEAK MATCH
NOT RECOMMENDED

18. AI Pipeline 6 — Job-Specific Resume



Resume Strategy AI

- Determine which experience should appear first.
- Identify the most important verified skills.
- Prioritize relevant achievements.
- Determine the appropriate summary emphasis.
- De-emphasize irrelevant content without deleting source facts.
- Identify legitimate keywords that can be used.

```
{
  "priority_skills": ["Cisco", "Network Security", "Firewall Management"],
  "priority_experience": ["ABC Technologies", "XYZ Networks"],
  "de_emphasize": ["Desktop Support"]
}
```

19. AI Pipeline 7 — Application Question Assistant

Draft role-specific answers using the candidate profile, job description, company context and the application question.

Sensitive declarations must never be guessed by AI, including:

- Work authorization.
- Sponsorship requirements.
- Disability disclosure.
- Veteran status.
- Criminal-history questions.
- Salary expectations when not previously set by the user.

20. Application Readiness Engine

Before submission, calculate a readiness state using deterministic and AI-validated inputs.

```
Job Match           88%
Resume Alignment    94%
Required Questions  100%
Required Documents  100%
Candidate Verification Complete

Result: APPLICATION READY
```

If any required information is uncertain, the state becomes USER ACTION REQUIRED.

21. Assisted Apply

Assisted Apply should be the default application mode.

```
AI Finds Job
↓
Match Analysis
↓
Resume Tailoring
↓
Application Preparation
↓
Candidate Review
↓
Browser Automation / Form Fill
↓
Candidate Submits
```

22. Controlled Auto Apply

Auto Apply may be added later, but should operate only inside strict user-defined rules and quality gates.

```
Minimum Match: 85%
Salary: $100,000+
Locations: Houston / Dallas / Remote
Employment: Full-time
```

Exclude: Contract / Staffing agencies
Maximum Applications: 10/day

Auto Apply Safety Gate

```
JOB FOUND
↓
Eligibility Rules
↓
Match Threshold
↓
Resume Generated
↓
Fact Validator
↓
Application Completion Check
↓
Policy Check
↓
SUBMIT or REVIEW REQUIRED
```

23. Application Tracker

Every application should have a persistent state.

```
Found
Reviewed
Skipped
Resume Preparing
Ready to Apply
Applied
Employer Viewed
Recruiter Contact
Interview
Second Interview
Offer
Rejected
Withdrawn
Hired
```

24. AI Pipeline 8 — Career Intelligence

Analyze outcome patterns across:

- Jobs viewed and matched.
- Jobs applied to.
- Resume versions.
- Recruiter response rate.
- Interview rate.
- Rejections.
- Offers.

Examples of useful insights:


- Which job titles produce the highest interview response.
- Which skills appear most often in missed opportunities.
- Which resume themes correlate with recruiter response.
- Whether the user is targeting jobs materially above or below demonstrated experience.

25. AI Job Coach

The AI Job Coach should use the user's actual Career Profile, application history and outcomes. It should not behave as an isolated generic chatbot.

- Why am I not getting interviews?
- Which roles appear to fit me best?
- Which missing skill appears most frequently?
- Which resume version performs better?
- What compensation ranges am I seeing?

26. AI Pipeline 9 — Interview Preparation




- Likely technical questions.
- Behavioral questions.
- Questions based on candidate gaps.
- STAR answer preparation.
- Company-specific preparation.
- Questions the candidate should ask.
- Future voice-based mock interview support.


27. AI Pipeline 10 — Job Quality and Fraud Analysis


Evaluate signals such as:

- Employer career-site verification.
- Duplicate job detection.
- Age of listing.
- ATS link validation.
- Suspicious domains or contact information.
- Unrealistic compensation.
- Repeated reposts.
- Third-party staffing reposts.



28. AI Pipeline Orchestrator





The orchestrator combines AI, deterministic rules, user data, job data, the Fact Ledger and application state. AI should not control workflow where deterministic logic is more reliable.

29. Three Types of Processing

A. Deterministic Processing

- Salary comparison.
- Dates and employment durations.
- Eligibility rules.
- Locations.
- Application states.
- Required fields.
- Scoring calculations.


B. AI Processing

- Resume interpretation.
- Professional rewriting.
- Semantic comparison.
- Job requirement interpretation.
- Interview preparation.

C. Hybrid Processing

- Job Match Score.
- Application Readiness.
- Resume Quality.
- Job authenticity assessment.

30. Multi-AI Review Pattern



Examples:

Resume:
Writer AI → Fact Reviewer AI → Fact Ledger Verification → Resume Output

Job Match:
Requirement Extractor → Matching Engine → Reasoning AI → Score Validator

31. AI Disagreement Handling

When independent AI systems materially disagree, route the task to a reconciliation stage rather than arbitrarily selecting one result.

```
AI 1 Assessment: Strong
AI 2 Assessment: Weak
    ↓
RECONCILIATION MODEL
    ↓
Original Job + Extracted Requirements + Candidate Facts + Conflicting Analyses
    ↓
Reconciled Structured Result
```

32. AI Confidence and Provenance

Structured AI outputs should carry confidence and source references where appropriate.

```
{
  "skill": "AWS",
  "candidate_has_skill": true,
  "confidence": 0.98,
  "source_fact_ids": [ "EXP-024", "SKILL-012" ]
}
```

Low-confidence claims should be suppressed or routed to user review.

33. AI Audit Log

```
{
  "task_id": "AI-82716",
  "user_id": "USR-101",
  "function": "resume_generation",
  "provider": "provider_a",
  "model": "configured_model",
  "input_version": "career_profile_v12",
  "output_version": "resume_v5",
  "status": "approved",
  "created_at": ""
}
```

Audit logs support debugging, hallucination investigation, model evaluation, cost analysis and support.

34. AI Cost Routing

Tier 1 — Low-Cost AI

- Classification.
- Extraction.
- Formatting.
- Normalization.

Tier 2 — General Intelligence

- Resume writing.
- Job analysis.
- Application response drafting.

Tier 3 — Advanced Reasoning

- Conflicting information.
- Resume validation.
- Complex matching.
- Reconciliation.
- Career analysis.

35. AI Failover

```
Primary Model
  ↓
Failure
  ↓
Retry
  ↓
Secondary Provider
  ↓
Fallback Model
```

Provider outages should degrade gracefully rather than disabling the entire platform.

36. Prompt Registry and Versioning

Prompts should be managed centrally and versioned.

```
CAREER_EXTRACTION_V1
JOB_EXTRACTION_V2
MATCH_REASONER_V3
RESUME_STRATEGY_V2
RESUME_WRITER_V4
RESUME_VALIDATOR_V3
APPLICATION_ANSWER_V2
INTERVIEW_PREP_V1
```

- Testable.
- Replaceable.
- Logged.
- Rollback-capable.
- Associated with model/version and output quality metrics.

37. Structured AI Output Contracts

Prefer JSON contracts between pipeline stages rather than passing unstructured prose from one model to another.

```
{
  "match_score": 88,
  "mandatory_requirements_met": 8,
  "mandatory_requirements_total": 9,
  "strengths": [],
  "gaps": [],
  "risk_flags": [],
  "recommended_resume_focus": []
}
```

38. User Dashboard

```
Good Morning, Michael

New Jobs Found           31
Strong Matches           8
Applications Ready       5
Need Your Review         3
Applications This Week   14
Interview Requests       2
```

Each job card should expose the user's match, resume status and next action.

39. Job Detail Page

- Company, title, location, salary and work arrangement.
- Source and detected/verified dates.
- Match score.
- Why the candidate matches.
- Missing requirements.
- Original resume alignment.
- Optimized resume alignment.
- Application readiness state.

40. Resume Version Control

The platform must retain the exact resume used for each application.

```
Master Resume
Resume V1
Resume V2
Microsoft Resume
Google Resume
Cisco Resume
```

This allows the candidate and interview-preparation engine to know exactly what was submitted.

41. Browser Extension

A browser extension may later:

- Detect job descriptions on supported sites.
- Import a job into the platform.
- Analyze candidate fit.
- Generate a tailored resume.
- Assist with form completion.
- Record the application.

```
AI JOB ASSISTANT

Match: 91%
Resume Optimized
14 application fields detected

[Prepare Application]
```

42. Job Discovery Strategy

Prioritize technically and contractually sustainable sources:

1. Employer career pages.
2. Public ATS feeds.
3. Licensed job APIs.
4. Employer partnerships.
5. Authorized aggregators.
6. Browser-assisted capture.

Do not make uncontrolled scraping the foundation of the business.

43. Privacy and User Control

- Users control personal and professional data.
- Sensitive information should be sent to AI providers only when necessary for a task.
- Work authorization and personal declarations should be explicitly confirmed.
- Resume uploads and application history require appropriate retention and deletion controls.
- AI outputs should be traceable to their source profile version.

44. MVP Development Scope — Phase 1

Candidate

- Resume-first onboarding and draft account creation.
- Email OTP / magic-link identity confirmation and consent capture.
- Dynamic missing-data interview based on extracted resume information.
- Registration and login.
- Resume upload.
- Master Career Profile.
- Fact Ledger.
- Resume Upscale.

Jobs

- Job ingestion.
- Normalization.
- Search.
- Recommendations.

AI

- Career Extractor AI.
- Career Reviewer AI.
- Job Requirement AI.
- Match AI.
- Resume Strategy AI.

- Resume Writer AI.
- Resume Validator AI.

Application

- Job-specific resume.
- Application Q&A.
- Review-and-apply workflow.
- Application tracker.

45. Phase 2

- Browser extension.
- Assisted autofill.
- Career Intelligence.
- Interview AI.
- Follow-up reminders.
- Job verification.
- Duplicate detection.
- Recruiter discovery.

46. Phase 3

- Controlled Auto Apply.
- AI Career Agent.
- Voice interview simulation.
- Salary negotiation assistant.
- Recruiter outreach assistant.
- Employer accounts.
- Candidate search.
- Employer job posting.

47. Subscription Model — Initial Product Hypothesis

Plan	Indicative Price	Primary Use	Examples
Free	\$0	Trial / entry	Career profile, basic resume, limited matching
Starter	\$7.99/mo	Resume improvement	Resume Upscale, more matches, tracking
Pro	\$14.99/mo	Active job seeker	Job-specific resumes, application assistance, analytics
Autopilot	\$24.99/mo	High automation	Controlled automation, advanced intelligence, interview AI

Final pricing should be validated against AI usage, job-data licensing, browser automation and support costs.

48. Core Product Differentiation


Do not market the product primarily as an AI Resume Builder.

Recommended positioning: YOUR AI JOB APPLICATION MANAGER

- Finds relevant opportunities.
- Understands requirements.
- Matches the candidate.
- Improves the resume.
- Prepares the application.
- Assists with submission.
- Tracks the process.
- Learns from outcomes.
- Prepares the candidate for interviews.

49. Key Architectural Principle


AI should operate on a structured, verified candidate record rather than repeatedly reinterpreting an unstructured resume.



Avoid the unreliable pattern:

Resume → AI → Random Rewrite

50. Recommended System Architecture



51. Development Rules to Lock

1. Never trust a single AI response for critical candidate facts.
2. Keep candidate facts separate from generated presentation.
3. Make AI providers replaceable.
4. Use different models for different functions.
5. Use deterministic code whenever AI is unnecessary.
6. Require structured JSON between AI stages.
7. Maintain provenance for generated information.
8. Never allow AI to fabricate candidate qualifications.
9. Keep humans in control of sensitive application answers.
10. Use independent AI reviewers for high-impact outputs.
11. Use expensive reasoning models only where they materially improve quality.
12. Maintain prompt and model versioning.
13. Log AI decisions for debugging and quality control.
14. Design Auto Apply around strict user-defined rules.
15. Optimize for high-quality relevant applications rather than maximum volume.

52. Product End State

The long-term platform should know:

- Who the candidate is.
- What the candidate has actually done.
- What the candidate can legitimately claim.
- Which jobs fit and which do not.
- How to position verified experience for each role.
- Where the candidate has applied.
- Which resume was used.
- Which employers responded.
- What patterns lead to interviews.
- Which skills are limiting opportunities.
- What action the candidate should take next.

End-state product: an AI Employment Operating System powered by Structured Candidate Intelligence + Job Intelligence + Multi-AI Reasoning + Application Automation + Outcome Learning.

53. Website & Landing Page UX / Conversion Design

Design objective

The public website should behave as the first step of the product rather than as a traditional marketing brochure. The primary conversion action is resume upload. The page should reduce friction, establish trust quickly, and move the visitor from curiosity to a useful job-search result with as few steps as possible.

Primary homepage structure

- Hero: split layout with product promise on the left and an immediate resume-upload action card on the right.
- Trust strip: concise privacy, security and control messaging directly below the hero.
- How it works: show the complete four-step flow from resume upload to job application.
- Right Job Fit: demonstrate intelligent matching and explain why the platform prioritizes quality over application volume.
- Tailored Application / Automation: show how the platform prepares and assists with applications using the verified candidate profile.
- Resume Upscale: explain how AI strengthens presentation without fabricating experience.
- Results / Efficiency: show measurable time savings, qualified applications and application outcomes.
- Testimonials / Social Proof: candidate stories focused on outcomes and saved effort.
- Pricing: simple, affordable plans with clear differences and low-friction entry.
- Final CTA: repeat the resume upload action rather than sending users to a generic registration form.

Hero section

Recommended desktop layout: approximately 55% message / 45% action card. On mobile, the upload card should move immediately below the headline and remain visible without excessive scrolling.

Left side content

- Clear outcome-driven headline, not a feature list.
- Short supporting sentence explaining that the platform builds the career profile, identifies strong-fit jobs, tailors applications and saves time.
- Optional small proof indicators such as applications prepared, time saved or strong-match jobs found once real data exists.
- Avoid generic stock-photo heroes that make the product look like a staffing agency or resume-writing service.

Right side action card

- Large drag-and-drop resume upload area.
- Accepted formats shown clearly.
- Primary CTA: Upload Resume / Start With My Resume.
- Small statement explaining that AI will build the user profile automatically.
- Privacy reassurance adjacent to the upload control.
- No lengthy signup form before upload.

Suggested hero messaging

Headline direction: Better applications. Less time. Lower cost.

Supporting direction: Upload your resume. We build your career profile, find the jobs you are best positioned for, tailor each application, and help you apply efficiently.

Primary CTA: Upload Resume

Secondary CTA: See How It Works

Trust and privacy strip

Because the first action asks for a resume containing personal data, trust must appear before or immediately after upload. Keep this short and visual.

- Your resume remains under your control.
- AI uses verified candidate facts rather than inventing credentials.
- Sensitive information is not submitted without user confirmation.
- Users can review or delete their profile and resume data.

How It Works section

Use a clean four-step visual flow instead of long explanatory text:

- 1. Upload Resume — AI creates the draft account, Fact Ledger and Master Career Profile.
- 2. Find Better Fits — platform analyzes job requirements against verified experience.
- 3. Tailor & Prepare — AI upscales and customizes the resume and application package.
- 4. Review & Apply — user approves, assisted automation fills repetitive application steps, and the application is tracked.

Right Job Fit section

This should be the first major product demonstration after onboarding. Show a real-looking product panel rather than abstract marketing graphics.

- Example job card with Match Score.
- Why You Match.
- Missing Requirements.
- Salary / location / work arrangement alignment.
- Strong Match / Good Match / Review labels based on configured thresholds.
- Message: the platform filters poor-fit roles so users spend time on opportunities worth pursuing.

Application Efficiency & Automation section

Present automation as controlled efficiency, not indiscriminate mass applying. The visual should show a job moving from matched → resume ready → application prepared → review → submitted.

- Application questions drafted from the Master Career Profile.
- Repeated form fields completed automatically where supported.
- Sensitive answers require confirmation.
- Application tracked after submission.
- Future controlled Auto Apply operates only inside explicit user rules.

Resume Upscale section

Resume Upscale should be visually differentiated from generic “AI rewrite” tools. Show before/after presentation while emphasizing that the underlying facts remain unchanged.

- Weak bullet identified.
- AI asks a targeted clarification question where necessary.
- Verified answer is added to the Fact Ledger.
- Resume statement is professionally rewritten.
- Final statement is validated against candidate facts.

Results and efficiency section

The product should eventually quantify value with user-specific metrics rather than vague claims.

Example weekly dashboard:

63 jobs analyzed | 18 strong matches | 12 applications submitted | 7 hours estimated time saved | 3 recruiter responses

Key positioning: More qualified applications per hour — not simply more applications.

Testimonials / social proof

Place testimonials after visitors have seen the product workflow. Testimonials should reinforce measurable value: better matches, time saved, stronger resumes, interviews and ease of use. Avoid generic quotes such as “Great website.”

Pricing presentation

Price is part of the differentiation, but the page should sell efficiency and outcomes first. Keep pricing simple and accessible, with a free entry point and clearly explained upgrade value. Do not position the brand primarily as “cheap”; position it as higher efficiency at a lower cost.

Visual direction

- Modern SaaS aesthetic with generous white space and strong typography.
- Use one primary brand color with a restrained accent system rather than excessive gradients.
- Use actual product UI mockups/cards as the dominant visual language.
- Use subtle motion for upload, match scoring and application progression.
- Avoid cluttered dashboards on the landing page; reveal one idea per section.
- Use diverse human imagery selectively in testimonials or outcome stories, not as the main hero visual.
- Design mobile-first for the upload and onboarding experience.

Homepage conversion flow

Traffic / Search / Social / Referral → Homepage Hero → Upload Resume → AI Builds Draft Profile → Identity & Consent Confirmation → Job Matches → First Tailored Application → Paid Conversion

Website conversion principle

Every major section should either increase trust, explain value, demonstrate product behavior, or return the visitor to the primary CTA. The homepage should not ask the visitor to “learn everything” before acting.

Updated positioning principle

The customer promise should combine affordability with efficiency: better-fit jobs, stronger truthful applications, less wasted time and a lower monthly cost. Price alone is easy to copy; efficiency and quality should remain the primary product differentiation.

54. Legal, Privacy, Consent & Data Governance Requirements

The legal and consent architecture is part of the product workflow and must be implemented in code, not treated only as footer documents. The platform should launch on the assumption of a U.S.-first, age-18+ consumer service unless a later jurisdiction-specific expansion is approved.

Resume upload is a consent event

Before a resume is transmitted, the upload area must state that uploading authorizes JobPilot to extract personal and professional information to create a temporary draft account and Career Profile and begin job-matching analysis. The user must confirm the resume is theirs or that they are authorized to upload it.

A resume upload creates a draft record only. The permanent account must not be activated until the user verifies an approved identifier and accepts the Terms of Service, Privacy Policy, and Resume Data & AI Processing Notice. Unactivated draft resume/profile data should be deleted or de-identified within 30 days, subject to limited fraud/security logs.

Data-use rules

- Do not sell resume, Career Profile, or application data for money or use it for cross-context behavioral advertising at launch.
- Do not allow third-party AI providers to use identifiable resume/application data to train publicly available or general-purpose models where contract/settings can prevent it.
- Collect and retain only data needed for the requested job-search service.
- Detect and suppress unnecessary high-risk data such as Social Security numbers, financial credentials, passwords, and unrelated medical information.
- Do not infer or use protected characteristics as job-match criteria.
- Voluntary EEO/demographic questions must not be answered automatically unless the user expressly elects to provide the answer.

Auto-Apply requires separate authorization

- Auto-Apply must be disabled by default and enabled only after a separate affirmative consent.
- The user must define eligibility rules, match threshold, exclusions, and an application cap.
- The system must stop for review on sensitive questions, legal attestations, background-check authorizations, contracts, medical/disability/EEO questions, and any field without a reliable verified answer.
- Disabling Auto-Apply must stop future automated submissions as soon as reasonably practicable.
- The platform must log the job, employer, date/time, resume version, material answers, and submission status for each automated application.

Subscription and billing consent

- Immediately before payment, show price, currency, billing frequency, renewal terms, trial conversion terms, and cancellation method.
- Recurring billing requires affirmative consent; no pre-checked boxes.

- Provide online cancellation through Account > Billing using a process at least as easy as enrollment.
- Maintain the version and timestamp of the billing terms accepted.

Required public legal pages

- Terms of Service
- Privacy Policy
- Resume Data Extraction, AI Processing & Account Creation Notice
- Auto-Apply Authorization & User Acknowledgment
- Subscription & Billing Terms
- Cookie Policy
- Acceptable Use Policy
- Job Listings & Third-Party Services Disclaimer

Required internal compliance standards

- Data Retention & Deletion Standard
- AI Governance & Data Use Standard
- Legal & Consent Implementation Checklist
- Vendor/subprocessor register
- Incident-response and breach-notification procedure

Privacy and user controls

- Account > Privacy should support profile correction, resume deletion, account deletion, export/access request, cookie settings, and application history.
- Users should be able to see why a job match was produced and correct profile facts that materially affect matching.
- If the service later introduces targeted advertising, sale/sharing of personal information, qualifying automated decisionmaking, or employer-side candidate screening, a new privacy/legal review is required before release.

Production security requirements

- Production resume upload must use HTTPS/TLS. Do not launch personal-data upload on plain HTTP.
- Encrypt sensitive resume/profile data at rest using appropriate platform controls.
- Require MFA for privileged/admin accounts and least-privilege access to user data.
- Do not write raw resumes, passwords, auth tokens, or sensitive application answers into general application logs.
- Propagate verified deletion requests to databases, object storage, vector stores, caches, search indexes, and subprocessors where applicable.

Legal-pack version

The initial JobPilot legal pack dated September 27, 2026 contains the public legal drafts and internal implementation standards. Counsel review is required before public launch and whenever product behavior materially changes.

Appendix A — Recommended Next Technical Specification

The next engineering document should convert this product specification into implementable backend contracts.

- Database schema and entity relationships.
- Service boundaries and API endpoints.
- AI Orchestrator state machine.
- Exact JSON input/output schemas for every AI stage.
- Prompt Registry definitions.
- Model-routing configuration.
- AI failover and retry logic.
- Validation and reconciliation rules.
- Job-source connector architecture.
- Application browser-extension architecture.
- Security, privacy and retention requirements.
- Telemetry, cost tracking and quality evaluation.
- Golden test fixtures for resume extraction, job matching and resume validation.

Legal Document Architecture - Version 2

Product decision: JobPilot will publish two primary public legal documents, with short contextual notices presented inside the product at the point of action.

Public document 1 - Terms of Service

- Account activation and service-use rules.
- Career-profile accuracy and no-fabrication rules.
- AI-assisted feature terms and limitations.
- Job-listing / third-party service disclaimer.
- Assisted Apply and Auto-Apply terms, including limited authorization and mandatory user-review categories.
- Subscriptions, automatic renewal, cancellation, refunds, and billing.
- Acceptable-use restrictions.
- IP, termination, warranty, liability, indemnification, dispute, and general legal terms.

Public document 2 - Privacy & Data Use Policy

- Resume extraction and temporary draft-account/profile creation.
- Categories and sources of personal information.
- AI processing, inferred data, model/provider controls, and model-training commitments.
- Job-search/application data and authorized employer/ATS transmission.
- Cookies, analytics, preference technologies, and cookie controls.
- Data sharing, retention, security, deletion, and user privacy rights.
- California, Texas, other state, and future international privacy requirements.

Just-in-time notices - UI only, not separate footer policies

- Resume Upload Notice: before the upload action, explain that uploading authorizes parsing/extraction to build a temporary profile; the permanent account still requires verification and acceptance of Terms/Privacy.
- Auto-Apply Enablement Notice: separate affirmative authorization before automated submissions are enabled; include the user-defined scope and categories that always require review.
- Recurring Billing Disclosure: immediately before purchase, show price, billing frequency, automatic renewal, trial conversion if any, and cancellation method.
- Cookie Settings: provide a persistent settings control where non-essential cookie consent or opt-out controls are required.

Footer/legal navigation

Recommended public footer links: Terms of Service | Privacy & Data Use | Cookie Settings (control, not a third policy). Additional state-specific opt-out links should appear only when required by the deployed data practices and applicable law.

Internal compliance documents

- Data Retention & Deletion Standard.
- AI Governance & Data Use Standard.
- Legal Consent Implementation Checklist.